

Stress-Testing Citation Faithfulness in Open-Weight Retrieval-Augmented Generation

Jeremiah O. Abimbola^{a,*}, Onyeka J. Nwobodo^c and Godlove Suila Kuaban^b

^a*Faculty of Automatic Control, Electronics and Computer Science, Silesian University of Technology, Gliwice, Poland*

^c*Department of Computer Graphics, Vision and Digital Systems, Faculty of Automatic Control, Electronics and Computer Science, Silesian University of Technology, Gliwice, Poland*

^b*Institute of Theoretical and Applied Informatics, Polish Academy of Sciences (IITIS-PAN), Bałtycka 5, Gliwice, Poland*

ARTICLE INFO

Keywords:

retrieval-augmented generation
citation faithfulness
attributed generation
mechanistic interpretability
probing

ABSTRACT

A document can support a generated claim without having influenced it. Extending prior citation-intervention work, we stress-test citation assignment in Llama-3.1-8B-Instruct, Qwen2.5-7B-Instruct, and Gemma-3-12B-it on 1,444 Natural Questions queries with KILT Wikipedia retrieval. We insert a target answer phrase into a random document, a relevant but originally uncited document, or a document cited elsewhere, then measure how often a recovered statement cites the altered document. Conditional adversarial-citation rates were 0–1.1% for random documents, reached 9.7% for relevant-uncited documents, and were 9.1%, 21.6%, and 17.8% for cited-other documents in Llama, Qwen, and Gemma. These are descriptive outcomes with different eligible and recovered denominators; document position and retention of the original source also differ across conditions. They do not isolate a relevance effect or reveal the source that caused an answer. Auxiliary analyses likewise do not validate provenance: ConflictBank answer matching is limited by predominantly ambiguous closed-book outputs, corrected hidden-state probes score below simple baselines, and selected component scaling corrects 0/7 missed and 0/58 spurious citations. The study provides a reproducible cross-model stress test of citation assignment and identifies the further controls needed before such behavior can support claims of faithful attribution.

1. Introduction

Retrieval-augmented generation (RAG) combines language models with retrieved documents Lewis, Perez, Piktus, Karpukhin, Goyal, Kuttler, Lewis, Yih, Rocktaschel, Riedel and Kiela (2020); Guu, Lee, Tung, Pasupat and Chang (2020); Izacard and Grave (2021). It is useful because models cannot store every relevant fact in their weights, their knowledge can become outdated, and they can produce plausible but unsupported answers Mallen, Asai, Zhong, Das, Khashabi and Hajishirzi (2023); Ji, Lee, Frieske, Yu, Su, Xu, Ishii, Bang, Madotto and Fung (2023). Retrieval supplies external documents that may be more current, specific, and easier to verify than stored knowledge. In practical RAG systems, citations connect an answer to its claimed evidence, allowing readers to inspect the source of a claim; they are therefore often treated as signs that the cited material supported the answer, particularly in information-seeking settings where users rely on citations to assess trustworthiness Liu, Zhang and Liang (2023).

A cited answer can look transparent while hiding several separate steps: the retriever selects documents, the generator writes an answer using both those documents and stored knowledge, and the system adds a citation marker pointing to one of the retrieved sources. Performing the final step well does not show that the earlier steps worked as intended. A system may retrieve a useful passage, ignore it, answer from memory, and then attach a citation that happens to match the answer, leaving readers with an output that appears grounded but does not explain how the answer was produced Wallat, Heuss, de Rijke and Anand (2025). Citations are therefore more than formatting; they serve as practical signals of provenance, auditability, and reliance on external evidence.

Producing citations is not the same as proving that they reflect what the model used. Attributed-generation systems have made progress in writing answers with citations and checking whether cited passages support individual claims Gao, Yen, Yu and Chen (2023); Liu et al. (2023); however, these checks mainly measure citation correctness or

*Corresponding author

✉ jeremiah.oluwagbemi.abimbola@polsl.pl (J.O. Abimbola); onyeka.nwobodo@polsl.pl (O.J. Nwobodo); gskuaban@iitis.pl (G.S. Kuaban)

ORCID(s): 0000-0002-9963-2563 (J.O. Abimbola)

completeness, asking whether a passage supports a claim and whether important claims have support, rather than whether the model used that passage instead of relying on stored knowledge or a shallow document-answer match. A citation can therefore be correct in content while still providing a misleading account of the model’s evidence use. For example, suppose a user asks, “What is the capital of Australia?” and the model already knows that the answer is Canberra. If the retrieved passage states that Canberra is the capital, the citation is correct, but the model may have produced the same answer even if that passage had been replaced with an unrelated document. The citation supports the claim, yet it does not by itself show that the passage influenced the answer.

We define *citation correctness* as whether a cited passage supports the claim made in the generated answer, whereas *citation faithfulness* asks whether that passage corresponds to evidence that actually influenced the model during generation. Correctness is a post-generation textual relation; faithfulness concerns evidence use and is therefore a causal question that requires testing how generation changes when the evidence changes. No single test answers this perfectly: behavioral interventions assess sensitivity to forged support, activation patching tests specific internal mechanisms, and probes reveal whether a risk signal is present in hidden states. Because these methods answer different questions, we treat them as complementary.

Wallat et al. formalize this distinction for RAG attributions and introduce post-rationalization interventions Wallat et al. (2025), in which an answer-bearing string is inserted into a document that was not the original factual source. When a model repeats the target statement and cites that document, the citation is difficult to interpret as true provenance. This approach connects with work that studies attribution from inside the model through context-sensitive token analysis and saliency-based document attribution Qi, Sarti, Fernández and Bisazza (2024), motivating evaluations that separate textual support from actual reliance on retrieved context.

Open-weight instruction-tuned models are useful for this study because they support controlled generation experiments across model families while providing access to hidden states and intermediate activations. We can therefore compare how models behave under document interventions, what information their representations contain, and whether selected internal components can change citation behavior Leon (2025). Access to activations does not, however, make faithfulness directly visible: a good probe demonstrates predictability, a successful patch demonstrates a causal effect under one intervention, and neither alone proves that a citation fully explains the model’s reasoning.

The primary contribution is a cross-model stress test of how citations follow answer-bearing insertions under three document-selection conditions. We extend the intervention approach of Wallat et al. Wallat et al. (2025) with explicit condition-availability and statement-recovery denominators. The comparison is descriptive: the conditions also differ in document position, source retention, and eligibility, so they cannot identify an isolated relevance effect or establish the provenance of a generated claim. ConflictBank, activation patching, and hidden-state probing are auxiliary diagnostics that test whether other measurements corroborate the behavioral interpretation. Their limited or negative results narrow that interpretation rather than supplying independent proof of a causal citation mechanism.

The rest of the paper is organized as follows. Section 2 reviews work on RAG, attributed generation, knowledge conflicts, mechanistic interpretability, and probing. Section 3 describes the datasets and model checkpoints, while Section 4 describes the research architecture and experimental methods. Section 5 presents the behavioral, conflict, mechanistic, and probe results. Section 6 discusses what the results mean for citation faithfulness. Sections 7 and 8 describe limitations and conclusions.

1.1. Research Objectives

This study asks how citation assignment changes when an answer phrase is inserted into a retrieved document. It uses an operational outcome—recovery of the target statement together with citation of the altered document—rather than claiming to observe which source actually caused the answer. The primary research question is: how often does this outcome occur across three document conditions and three open-weight models, after accounting for intervention availability and statement recovery? Two auxiliary questions ask whether knowledge-conflict answers and citations provide supporting context, and whether hidden-state probes or activation interventions offer independent predictive or causal evidence. These auxiliary analyses cannot validate the primary attribution claim on their own.

2. Related Work

Early RAG systems joined language models with external retrieval for knowledge-intensive tasks. Dense Passage Retrieval learns question and passage representations with a dual encoder, providing a learned alternative to sparse lexical retrieval for open-domain question answering Karpukhin, Oguz, Min, Lewis, Wu, Edunov, Chen and Yih

(2020). REALM adds retrieval during language-model pre-training. RAG generates while marginalizing over retrieved passages. Fusion-in-Decoder lets the decoder combine evidence from several passages Guu et al. (2020); Lewis et al. (2020); Izacard and Grave (2021). Self-RAG extends this line of work by training a model to decide when to retrieve and to critique its own use of retrieved passages, improving factuality without requiring a fixed retrieval schedule Asai, Wu, Wang, Sil and Hajishirzi (2024); even so, a self-critique signal is not itself a test of whether a specific citation reflects the passage that drove a specific claim. These systems showed that retrieval can improve factual and open-domain question answering. However, the usual metrics mainly measure answer quality. Natural Questions provides real user questions with annotated answers, and KILT gives a shared Wikipedia source with provenance annotations Kwiatkowski, Palomaki, Redfield, Collins, Parikh, Alberti, Epstein, Polosukhin, Devlin, Lee, Toutanova, Jones, Kelcey, Chang, Dai, Uszkoreit, Le and Petrov (2019); Petroni, Piktus, Fan, Lewis, Yazdani, De Cao, Thorne, Jernite, Karpukhin, Maillard, Plachouras, Rocktaschel and Riedel (2021). These datasets make retrieval and source tracking easier to evaluate. Still, a relevant passage and a correct answer do not prove that the generator relied on that passage for a specific claim.

Attributed generation makes sources visible by asking systems to attach references to generated claims. Gao et al. compare inline and post-hoc citation generation, and they evaluate citation quality through claim-level support and completeness Gao et al. (2023). Liu et al. study generative search engines from the reader's point of view. They measure whether claims are supported by citations and whether users can verify those claims Liu et al. (2023). Rashkin et al. formalize a closely related comparison with the Attributable to Identified Sources (AIS) framework, which asks whether a generated statement is verifiable against an independent source regardless of the generation task Rashkin, Nikolaev, Lamm, Aroyo, Collins, Das, Petrov, Tomar, Turc and Reitter (2023), and Dziri et al. extend AIS-style annotation to knowledge-grounded dialogue with the BEGIN benchmark, finding that automatic attribution metrics often rely on spurious correlations rather than genuine source use Dziri, Rashkin, Linzen and Reitter (2022). This work gives important standards for useful citations: important claims should be covered, cited sources should support them, and references should be easy to inspect. But these standards mainly compare the final claim with the final source. They do not test whether the source helped produce the claim. A model can answer from memory and then attach a passage that supports the answer. That citation may be correct, but it may still give a misleading account of provenance.

RAGTruth provides complementary evaluation data with response-level and word-level annotations of hallucinations in retrieval-augmented outputs Niu, Wu, Zhu, Xu, Shum, Zhong, Song and Zhang (2024). Such annotations identify unsupported or contradictory content; citation-faithfulness interventions additionally ask whether the cited evidence influenced generation.

Citation faithfulness asks whether an emitted citation reflects the evidence that influenced generation. Wallat et al. define the difference between correctness and faithfulness in RAG and introduce adversarial interventions Wallat et al. (2025). Their method places answer-bearing strings in documents that were not the original source of support. If a model repeats the target statement and cites that inserted document, this is an operational adversarial-citation event; it does not by itself identify the source that caused the answer. Qi et al. study faithful attribution from inside the model by using context sensitivity and saliency to estimate which retrieved documents affect answer tokens Qi et al. (2024). These approaches answer different questions. Adversarial tests show behavior under controlled input changes. Internal attribution methods estimate influence in one forward pass.

This distinction follows the broader interpretability requirement that explanations reflect model behavior, rather than merely appear plausible to a reader Jacovi and Goldberg (2020). The attention-as-explanation debate illustrates why that requirement needs explicit tests: Jain and Wallace find that substantially different attention distributions can yield equivalent predictions Jain and Wallace (2019), while Wiegreffe and Pinter challenge the generality of that conclusion and propose additional model-aware diagnostics Wiegreffe and Pinter (2019). These studies motivate testing attribution claims under interventions rather than treating a visible weight or citation marker as sufficient evidence of faithful explanation.

The gap between what a model outputs and what actually drove that output is documented outside citations as well. Abstractive summarizers frequently hallucinate content that is unsupported by the source document despite reading fluently Maynez, Narayan, Bohnet and McDonald (2020), and chain-of-thought explanations can rationalize an answer without reflecting the computation that produced it, so that a model's stated reasoning and its true reasoning can diverge Turpin, Michael, Perez and Bowman (2023). Citation post-rationalization is a structurally similar failure mode applied to retrieval evidence rather than to a summary or a reasoning trace, which situates the present study within a broader pattern of unfaithful post-hoc explanation in generative models.

The difference between retrieved evidence and stored model knowledge also appears in work on knowledge conflicts. Petroni et al. show that pretrained language models can recall relational facts through cloze queries without task-specific fine-tuning, establishing a basis for studying factual knowledge stored in model parameters Petroni, Rocktäschel, Riedel, Lewis, Bakhtin, Wu and Miller (2019). PopQA studies when models answer from stored knowledge and when outside evidence helps, with entity popularity as an important factor Mallen et al. (2023). ConflictBank tests cases where supplied context conflicts with a model’s stored knowledge and measures which source the model follows Su, Zhang, Qu, Zhu, Li, Sun, Li, Zhang and Cheng (2024). This line of work builds on earlier studies of parametric-contextual conflict: Longpre et al. formalize entity-substitution knowledge conflicts and show that models often default to memorized answers even when substituted context is provided Longpre, Perisetla, Chen, Ramesh, DuBois and Singh (2021), and Neeman et al. train models to output separate parametric and contextual answers so the two knowledge sources can be inspected independently Neeman, Aharoni, Honovich, Choshen, Szpektor and Abend (2023). Xu et al. survey this broader area, distinguishing context-memory, inter-context, and intra-memory conflicts Xu, Qi, Guo, Wang, Wang, Zhang and Xu (2024). Related work on parametric knowledge itself, such as TruthfulQA, shows that a model’s stored answer is not always trustworthy either, since larger models can be more prone to reproduce common misconceptions Lin, Hilton and Evans (2022), which further motivates testing whether a citation reflects genuine reliance on supplied evidence rather than defaulting to either source. Following the supplied context is related to citation faithfulness, but it is not the same thing. A context-matching answer shows that the model is sensitive to the supplied evidence. A faithful citation must also point to the document that actually supported the claim.

Mechanistic interpretability gives tools for testing causal ideas inside a model. Geva et al. interpret transformer feed-forward layers as key-value memories whose keys match textual patterns and whose values induce distributions over the output vocabulary, providing a functional account of components examined in internal attribution studies Geva, Schuster, Berant and Levy (2021). Causal tracing and activation editing have found transformer components involved in factual recall, and they show that changing internal states can change model predictions Meng, Bau, Andonian and Belinkov (2022). TransformerLens provides hooks and activation caches that are commonly used for these patching experiments Nanda and Bloom (2022). Beyond single-fact factual recall, activation patching has also been used to reverse-engineer complete multi-component circuits for a specific behavior, such as the indirect-object-identification circuit in GPT-2 small, which spans 26 attention heads across several layers Wang, Variengien, Conmy, Shlegeris and Steinhardt (2023); this illustrates that a behavior can depend on a distributed combination of components rather than on any single one, a possibility that our limited mechanistic analysis cannot establish. The closest related work is by van Dort and Heuss, who study citation generation in Llama-3.1-8B-Instruct with activation patching on PopQA and a citation-token logit-difference metric van Dort and Heuss (2026). Their work shows how activation patching can be used to study citation-token behavior in a controlled setting. Our mechanistic experiment should be read as a replication-and-extension attempt rather than a new localization method: it applies a related citation-logit setup to a small PopQA capital-city sample, while the behavioral benchmark, conflict analysis, and hidden-state probes extend the evaluation beyond citation-token logits.

Linear probes ask a different question: can a simple classifier recover a target property from a model representation? Prior work has used such methods to study whether language-model activations contain information about truth and uncertainty. Burns et al. recover latent answers to true-or-false questions by finding directions in activation space that satisfy logical consistency constraints, even without labeled examples Burns, Ye, Klein and Steinhardt (2023). Azaria and Mitchell train classifiers on hidden activations to distinguish true from false statements, including statements generated by the model itself Azaria and Mitchell (2023). Li et al. train linear probes on attention-head activations to locate directions associated with truthful answers and then use those directions for inference-time intervention Li, Patel, Viégas, Pfister and Wattenberg (2023a).

These studies show that internal representations can support prediction of factual truth, latent knowledge, or semantic uncertainty. Our probe target is different: it is not whether an answer is factually correct or whether the model is uncertain, but whether a citation is post-rationalized under a controlled document intervention. This places citation faithfulness within the broader use of hidden-state probes for model monitoring while introducing an attribution-specific behavioral label. As with probing more generally, strong classification performance shows that information linked to the label is linearly available; it does not prove that the model uses the probed direction or that the direction causes the behavior Belinkov (2022). Probe evidence must therefore be interpreted separately from causal intervention evidence.

This distinction matters increasingly to information and computer science audiences. Model-internal attribution has been evaluated with context sensitivity and saliency Qi et al. (2024); knowledge-conflict benchmarks now quantify how supplied evidence competes with parametric memory Su et al. (2024); and mechanistic work examines citation-token behavior directly Wallat et al. (2025); van Dort and Heuss (2026). Our contribution is a cross-model, three-condition descriptive stress test with explicit recovery and eligibility denominators. The question sample and answer-string intervention build on Wallat et al. (2025); the auxiliary conflict, patching, and probe analyses test the limits of stronger interpretations rather than constituting four independent demonstrations of citation faithfulness.

Complementary studies reinforce the importance of evaluating retrieval and evidence use as separate parts of an information system. Li et al. formulate conversational question answering as generative retrieval, showing that passage selection is itself a central modeling problem rather than a transparent precondition for answer generation Li, Yang, Wang, Wei and Li (2023b). Liu et al. likewise show that query generation and the management of conversational context affect legal-case retrieval, illustrating how changes in the information supplied to a model can alter the evidence available for an answer Liu, Wu, Zhang, Liu, Wang, Li, Zhang and Ma (2022). For evaluation, Wang et al. develop a scalable cross-language fact-checking study over 4,500 claims in nine languages and report substantial variation across languages and models Wang, Tsang and Zhou (2026). Together, these studies support a broader evaluation perspective: retrieval quality, contextual evidence selection, answer correctness, and the provenance of generated claims should be measured separately. Our study adds a document-insertion stress test of citation assignment and reports where its auxiliary internal and knowledge-conflict diagnostics fail to identify provenance.

3. Datasets

The behavioral sample contains 1,444 unique Natural Questions identifiers and questions released with the experiments of Wallat et al. (2025). The questions originate from Natural Questions Kwiatkowski et al. (2019); the retrieved evidence comes from the full KILT Wikipedia knowledge source Petroni et al. (2021). We converted Wikipedia pages into non-overlapping chunks of 100 whitespace tokens before BM25 retrieval and retained the top five chunks per query. Thus, Natural Questions defines the question sample and KILT defines the evidence population; the intervention counts in Table 2 are lower because they are conditional on an original cited statement and a constructible adversarial document.

The mechanistic sample came from the `akariasai/PopQA` test split Mallen et al. (2023). We retained capital-city relations and selected 31 clean/corrupted pairs using the citation-logit filter described below. The counterfactual sample came from the `Warrieryes/CB_qa` training split associated with ConflictBank Su et al. (2024). We streamed that split, retained the first 2,000 examples with the expected schema, and sorted them by example identifier before generation; therefore, the ConflictBank results contain 2,000 examples per model. ConflictBank was used only for an auxiliary knowledge-conflict analysis, PopQA only for a mechanistic diagnostic, and probe labels and features came from our behavioral interventions rather than either dataset.

All model weights and tokenizers came from their official Hugging Face releases. We used three instruction-tuned checkpoints: `meta-llama/Llama-3.1-8B-Instruct`, `Qwen/Qwen2.5-7B-Instruct`, and `google/gemma-3-12b-it`. These models were released by Meta, Alibaba’s Qwen team, and Google DeepMind, respectively Meta AI (2024); Qwen Team (2024); Google DeepMind (2025). Llama 3.1 and Gemma 3 require license acceptance. Qwen2.5 is distributed under the Apache 2.0 license. Because these are open-weight models, we can generate answers, extract hidden states, and run activation-patching experiments. For the semantic-similarity baseline, we used `sentence-transformers/all-mpnet-base-v2` to encode document-query pairs, following the sentence-transformer bi-encoder approach Reimers and Gurevych (2019).

4. Methodology

4.1. Research Architecture

The primary path uses Natural Questions questions and top-5 BM25 retrieval from 100-token KILT Wikipedia chunks. It runs deterministic generation in three models, inserts answer phrases into selected documents, and measures recovered statements that cite the altered document. Hidden-state probes use labels from this path. Separate auxiliary paths use ConflictBank to examine context–answer matching and PopQA capital-city pairs for Llama-only activation patching. Figure 1 summarizes these distinct roles.

Dataset/source	Sample or coverage	Experimental role
Natural Questions Kwiatkowski et al. (2019)	1,444 unique questions	Main behavioral question sample
KILT Wikipedia Petroni et al. (2021)	Full knowledge source; 100-token chunks and top-5 retrieval	Retrieved evidence for the behavioral benchmark
PopQA Mallen et al. (2023)	31 clean/corrupted capital-city pairs	Llama-only mechanistic activation patching
ConflictBank Su et al. (2024)	2,000 examples per model	Auxiliary context-answer matching analysis

Table 1

Summary of datasets and their roles in the citation-faithfulness experiments.




Figure 1: Research architecture for the citation-faithfulness experiments. The main path evaluates retrieval, generation, behavioral interventions, and probes across all three models. The auxiliary mechanistic path performs Llama-only activation patching on PopQA examples.

4.2. Task Setup

Each example begins with a question and a document collection. Let q_i be the question for example i , and let $D = \{d_1, \dots, d_N\}$ be the document corpus. The retriever ranks documents with a BM25 score $s(q_i, d)$ Robertson and Zaragoza (2009) and returns the top- k documents:

$$R_k(q_i) = \text{TopK}_{d \in D} s(q_i, d), \quad (1)$$

In this study, $k = 5$. We write the retrieved set as $C_i = \{d_{i1}, \dots, d_{ik}\}$. The model receives the question and these retrieved documents as a prompt $x_i = (q_i, C_i)$ and generates an answer:

$$y_i = M_\theta(x_i), \quad (2)$$

Here, M_θ is one of the instruction-tuned models. The generated answer contains one or more statements and citation markers. For a target statement t_i and a cited document index c_i , citation correctness can be written as a support relation:

$$\text{Correct}(t_i, c_i) = \mathbb{1}\{d_{ic_i} \models t_i\}, \quad (3)$$

where $d_{ic_i} \models t_i$ means that the cited document supports the statement. This only checks the final text. It does not show whether the model used that document while producing the answer.

To stress-test citation assignment, we change the retrieved context and observe the model's response. For each target statement, we create an intervention I_a that inserts an answer-bearing phrase into an adversarial document d_{ia} . The new retrieved context is

$$C_i^{(a)} = I_a(C_i), \quad (4)$$

and the model generates a new answer

$$y_i^{(a)} = M_\theta(q_i, C_i^{(a)}). \quad (5)$$

We then define two indicators. The first checks whether the target statement is recovered:

$$r_i^{(a)} = \mathbb{1}\{t_i \subset y_i^{(a)}\}. \quad (6)$$

The second checks whether the adversarial citation marker is used in the sentence that contains the recovered target phrase:

$$z_i^{(a)} = \mathbb{1}\{[a] \in \text{sent}(y_i^{(a)}, t_i)\}. \quad (7)$$

The post-rationalization label is therefore

$$\ell_i^{(a)} = r_i^{(a)} z_i^{(a)}. \quad (8)$$

A value of $\ell_i^{(a)} = 1$ means that the model recovered the target statement and cited the adversarial document. This operational event is the main behavioral outcome. It shows citation co-occurrence after insertion; it does not establish that the inserted document caused the answer or that the citation is unfaithful.

4.3. Models and Execution

We evaluated three open-weight instruction-tuned models: meta-llama/Llama-3.1-8B-Instruct, Qwen/Qwen2.5-7B-Instruct, and google/gemma-3-12b-it. Behavioral and probe experiments were run for all three models. Mechanistic activation patching was run only for Llama-3.1-8B-Instruct, which matches the mechanistic scope of the study. Generation used greedy decoding with a maximum of 256 new tokens. The experiments ran on Apple MPS with backend-selected reduced precision and no quantization.

We fixed the seed at 42 and used each model's official chat template, `do_sample=false`, one beam, and 256 new tokens. The system prompt told the model to use relevant documents, cite factual statements immediately with document numbers [1]–[5], cite only supporting documents, avoid citing a document merely because it shared an answer phrase, and answer from model knowledge without a citation when the documents lacked the answer. Each user prompt contained five numbered documents and one question. We saved the rendered prompt, input length, raw output, parsed citations, and a SHA-256 prompt hash for each generation.

For each original answer, we chose one statement with a single citation and a target phrase of at most eight tokens that appeared in its cited document. We tested three conditions: a random document that excluded the original pages and target phrase, a relevant retrieved document that the original answer did not cite, and another document cited elsewhere in that answer. We inserted the target phrase into the selected document. We labeled an output POST-RATIONALIZED only when it recovered the normalized target phrase and cited the altered document in the same statement.

4.4. Behavioral Post-Rationalization Benchmark

The main behavioral benchmark uses 1,444 Natural Questions queries and top-5 BM25 retrieval from KILT Wikipedia. For each model, we first generated answers with citations. The system instruction required citations for document-supported claims. We then parsed the cited statements and built interventions around target statements that the model originally cited.

Each target statement was tested under three adversarial conditions:

- **Random:** the answer-bearing phrase is inserted into an otherwise random document, placed at citation index [5].
- **Relevant uncited:** the phrase is inserted into a retrieved document that is topically relevant but was not the original cited source.
- **Cited other:** the phrase is inserted into another document that the model cited elsewhere in the original answer.

An intervention is available when we can construct the adversarial document. We label the output as POST-RATIONALIZED when two things happen: the model recovers the target statement, and the adversarial document's citation marker appears in the sentence with the recovered target phrase. The main metric is the adversarial citation rate among available interventions that recover the target statement:

$$\text{PRR} = \frac{\#\{\text{recovered target and adversarial citation}\}}{\#\{\text{condition available and recovered target}\}}. \quad (9)$$

Wilson 95% intervals were computed for the conditional rate.

The original supporting document remains in the relevant-uncited and cited-other contexts. The random condition replaces document [5]; therefore, it also replaces the original supporting document when that document occupied slot [5] (18 Llama cases, 37 Qwen cases, and 113 Gemma cases). This position and source-retention difference limits the random condition as a strict control.

As a concrete example, Qwen2.5-7B-Instruct was asked, "When did Madras change its name to Chennai?" In the original run, it answered "Madras changed its name to Chennai in 1996 [5]." For the relevant-uncited intervention, the phrase "1996" was appended to retrieved document [1], a relevant Madras Railway passage that was not the original cited source. The model then produced "Madras changed its name to Chennai in 1996 [1]." Because the target statement was recovered and the newly modified document was cited in the same sentence, this example receives the label POST-RATIONALIZED. It illustrates an operational citation shift: both documents can support the statement, and the altered document was cited after the answer phrase was placed in it. This single example does not establish the model's internal causal process; the benchmark measures the observable behavior systematically across many examples.

4.5. Auxiliary Knowledge-Conflict Analysis

We use ConflictBank as an auxiliary context-sensitivity analysis, not as validation of citation provenance or as probe training data. ConflictBank examples contain a question, an original or parametric answer, a conflicting answer, and a piece of evidence that supports the conflicting answer. This lets us test what the model does when the retrieved context disagrees with what the model may already know.

For each example j , let q_j be the question, a_j^p be the original or parametric answer, a_j^c be the conflicting answer, and e_j^c be the conflicting evidence. Each model is run twice. First, it answers the question without any supporting document:

$$y_j^0 = M_\theta(q_j). \quad (10)$$

Second, it answers the same question with the conflicting evidence provided as the retrieved document:

$$y_j^c = M_\theta(q_j, e_j^c). \quad (11)$$

The generated answer is then assigned to one of three classes. After text normalization, it is labeled PARAMETRIC MATCH if it contains a_j^p but not a_j^c . It is labeled CONTEXT MATCH if it contains a_j^c but not a_j^p . It is labeled AMBIGUOUS if it contains both answers or neither answer:

$$g(y_j) = \begin{cases} \text{PARAMETRIC MATCH,} & a_j^p \subset y_j \text{ and } a_j^c \not\subset y_j, \\ \text{CONTEXT MATCH,} & a_j^c \subset y_j \text{ and } a_j^p \not\subset y_j, \\ \text{AMBIGUOUS,} & \text{otherwise.} \end{cases} \quad (12)$$

We apply this classifier to both y_j^0 and y_j^c . The descriptive signal is the rate at which the with-context answer is labeled CONTEXT MATCH:

$$\text{CMR} = \frac{1}{n} \sum_{j=1}^n \mathbb{1}\{g(y_j^c) = \text{CONTEXT MATCH}\}. \quad (13)$$

In this study, $n = 2000$ for each model. A high context-match rate shows that the output often contains the supplied answer under this matcher; without a recognizable closed-book answer, it does not measure a shift away from stored knowledge. The resulting class counts are converted to percentages for the context-shift figure in Section 5.

Because the no-context answers are dominated by the AMBIGUOUS class, the unconditional context-match rate cannot by itself quantify a shift *away from parametric knowledge*: for most examples we have no evidence that the model knew the fact in the first place. We therefore report two additional measures that condition on demonstrated knowledge and connect the conflict to the citation. First, restricting to examples the model answered correctly without context ($g(y_j^0) = \text{PARAMETRIC MATCH}$), the conditional override rate is

$$\text{CMR}_{|\text{known}} = \frac{\#\{g(y_j^c) = \text{CONTEXT MATCH and } g(y_j^0) = \text{PARAMETRIC MATCH}\}}{\#\{g(y_j^0) = \text{PARAMETRIC MATCH}\}}. \quad (14)$$

Second, among with-context answers labeled CONTEXT MATCH, we record whether the answer cites the supplied conflicting document [1], giving a context-citation co-occurrence rate. Because the prompt requests citation of supplied evidence, this rate cannot determine which source caused the claim. The plain substring matcher used for classification is replaced by a content-token matcher that treats a candidate answer as present when its stopword-filtered tokens appear as a contiguous run in the normalized output, so that verbose parametric answers are not spuriously labeled ambiguous. Both measures are recomputed from the saved generations without additional inference.

4.6. Mechanistic Activation Patching

For mechanistic analysis, we built capital-city pairs from PopQA [Mallen et al. \(2023\)](#). A clean prompt contains a document of the form “A is the capital of S” and asks the matching question. A corrupted prompt replaces the clean document with a distractor document. The distractor is length-matched in tokenization for the subject and answer spans. We selected the first 31 examples where the citation logit difference is positive in the clean prompt and negative in the corrupted prompt:

$$\Delta = \text{logit}(\text{" ["}) - \text{logit}(\text{" . "}). \quad (15)$$

For each selected example, we cached clean and corrupted activations. We then patched residual streams, MLP outputs, and attention heads. Denoising recovery and noising degradation were computed as

$$R_{\text{denoise}} = \frac{\Delta_{\text{patched}} - \Delta_{\text{corrupt}}}{\Delta_{\text{clean}} - \Delta_{\text{corrupt}}}, \quad D_{\text{noise}} = \frac{\Delta_{\text{clean}} - \Delta_{\text{patched}}}{\Delta_{\text{clean}} - \Delta_{\text{corrupt}}}. \quad (16)$$

We marked components as positive if their mean denoising recovery was at least 0.10. We marked components as necessary if their mean noising degradation was at least 0.20. We then tested the selected components by scaling them during generation on missed-citation and spurious-citation examples. To assess whether a null selection reflects low power at the threshold rather than a genuine absence of effect, we additionally bootstrapped the mean denoising recovery of the best-scoring component of each type (attention head, MLP write, and residual position) over the selected examples and report percentile 95% intervals.

4.7. Linear Faithfulness Probes

For each model, we extracted hidden-state features from behavioral intervention examples labeled POST-RATIONALIZED or NOT-POST-RATIONALIZED. To prevent the feature location from depending on the label, we anchored every example at the same structural position: the final token of the target phrase, which is the point at which the model

must decide whether to attach a citation. This anchor is chosen identically for both classes and precedes any citation marker, removing the identified label-dependent anchor. Other output, lexical, or condition cues may still remain in the features. For each layer, the feature vector was the mean hidden state over the anchor token and the two answer tokens before it.

For layer λ and example i , let $h_{i,\lambda,t}$ be the hidden state at answer-token position t . If T_i is the anchor token position, the probe feature is

$$u_{i,\lambda} = \frac{1}{3} \sum_{t=T_i-2}^{T_i} h_{i,\lambda,t}. \quad (17)$$

For each layer, the probe predicts the probability of post-rationalization with logistic regression:

$$\hat{p}_{i,\lambda} = \sigma(w_\lambda^\top u_{i,\lambda} + b_\lambda), \quad (18)$$

where σ is the sigmoid function. The probe is trained with balanced binary cross-entropy,

$$\mathcal{L}_\lambda = - \sum_i \alpha_{\ell_i} [\ell_i \log \hat{p}_{i,\lambda} + (1 - \ell_i) \log(1 - \hat{p}_{i,\lambda})], \quad (19)$$

where ℓ_i is the post-rationalization label and α_{ℓ_i} is the class weight. This weighting is used because post-rationalized examples are much rarer than non-post-rationalized examples.

We split questions by `question_id` into 70% training, 15% validation, and 15% test sets. This avoids leakage across interventions for the same question. We trained one L2 logistic regression probe per layer with balanced class weights. We selected the layer with the highest validation AUROC and report test AUROC, AUPRC, and balanced accuracy in Table 5. To test whether the probe captures a condition-general faithfulness signal rather than cues specific to one intervention type, we additionally performed leave-one-condition-out transfer at the selected layer: for each intervention condition we trained on the remaining conditions and evaluated on the held-out condition, reporting transfer AUROC and AUPRC in Section 5. The baselines used lexical overlap, citation-logit margin, and document-query cosine similarity from a sentence-transformer encoder. Lexical overlap was the Jaccard similarity between normalized target-statement and adversarial-document tokens; the citation-logit baseline was the marker-versus-period logit margin; and the semantic baseline was cosine similarity between question and adversarial-document embeddings from `sentence-transformers/all-mpnet-base-v2`. Baseline scores were evaluated on the same question-level test split as the selected-layer probes.

5. Experimental Results

5.1. Behavioral Citation Shifts

The behavioral benchmark compares three types of forged documents: random documents, relevant documents that were not originally cited, and documents cited elsewhere in the original answer. The observed conditional rates are lowest in the random condition and higher in the relevant-uncited and cited-other conditions. Table 2 reports the eligible and recovered denominators. These are descriptive comparisons because document position, source retention, and eligibility differ by condition.

The rates are conditional on different funnels rather than a common denominator. The original-answer files contain 1,444 generations per model; parseable cited answers numbered 629 for Llama, 916 for Qwen, and 1,106 for Gemma. Target extraction then yielded 274, 326, and 829 question-level intervention sets, respectively. Available cited-other cases were only 23, 59, and 117, so the cited-other percentages should not be treated as a definitive cross-model ranking.

5.2. Auxiliary Knowledge-Conflict Analysis

The ConflictBank experiment compares each model’s answer without supporting context with its answer after conflicting evidence is supplied. Each model therefore has two bars in Figure 2, with stacked colors showing whether the answer matches the model’s parametric answer, matches the supplied context, or remains ambiguous under the

Model	Condition	Events	Available	Recovered	Rate	Wilson 95% CI
Llama-3.1-8B	Random	0	274	226	0.0%	[0.0, 1.7]
Llama-3.1-8B	Relevant uncited	8	245	231	3.5%	[1.8, 6.7]
Llama-3.1-8B	Cited other	2	23	22	9.1%	[2.5, 27.8]
Qwen2.5-7B	Random	3	326	274	1.1%	[0.4, 3.2]
Qwen2.5-7B	Relevant uncited	27	299	279	9.7%	[6.7, 13.7]
Qwen2.5-7B	Cited other	11	59	51	21.6%	[12.5, 34.6]
Gemma-3-12B	Random	4	829	682	0.6%	[0.2, 1.5]
Gemma-3-12B	Relevant uncited	13	726	681	1.9%	[1.1, 3.2]
Gemma-3-12B	Cited other	18	117	101	17.8%	[11.6, 26.4]

Table 2

Operational adversarial-citation events. Available counts are constructible interventions; recovered counts are the denominators for rates and Wilson intervals.

string-matching rule. We use this figure instead of a separate table because the main result is the change between the two settings. Without context, most outputs are ambiguous because the short matching rule often finds neither answer directly; with conflicting evidence, all three models shift toward the context-matching class. This is evidence of context sensitivity under the classifier, not a clean before–after measure of overriding parametric knowledge: the no-context condition was ambiguous for 1,994 Llama, 1,942 Qwen, and 1,916 Gemma examples. We therefore interpret the figure as a context-conditioned answer-class analysis rather than a definitive measure of knowledge conflict resolution.




Figure 2: ConflictBank answer classes before and after conflicting evidence is supplied. The figure shows answer classes under the content-token matcher; most closed-book outputs are ambiguous, so the bars do not quantify a general shift from parametric knowledge.

Table 3 reports the two knowledge-conditioned measures. Restricting to the examples that each model answered correctly without context isolates a small but interpretable subset (56 Qwen, 71 Gemma, and 5 Llama). Within this subset, the supplied conflicting evidence overrides the model’s own prior answer in 78.6% of Qwen, 67.6% of Gemma, and 80.0% of Llama cases, so the shift is measured only where a parametric answer was actually demonstrated. The Llama subset is too small to support a precise estimate and is reported for completeness. Among with-context answers matching the supplied answer, the conflicting document was cited in 94.6%, 98.6%, and 96.4% of Qwen, Gemma, and Llama cases. This is answer–citation co-occurrence under an instruction to cite supplied evidence, not a test of whether that document caused the generated claim.

Model	Parametric-known	Conditional override (Wilson 95%)	Context-citation rate
Llama-3.1-8B	5	80.0% [37.6, 96.4]	96.4%
Qwen2.5-7B	56	78.6% [66.2, 87.3]	94.6%
Gemma-3-12B	71	67.6% [56.1, 77.3]	98.6%

Table 3

Knowledge-conditioned ConflictBank measures. The conditional override rate is computed only on examples answered correctly without context; the context-citation rate is the share of context-matching answers that cite the supplied conflicting document.

Validation set	Count	Correction success rate
Missed citation	7	0.0%
Spurious citation	58	0.0%

Table 4

Mechanistic validation after component selection and scaling. Correction success rate is the percentage of examples whose citation behavior is fixed by the scaling intervention.

5.3. Activation-Patching Diagnostic

The PopQA mechanistic phase selected 31 clean/corrupted pairs that met the filtering rule. Patching produced 512 residual entries, 512 MLP entries, and 1,024 attention-head entries. Component selection found no positive citation heads or MLP components under the denoising threshold. It found one necessary attention head, layer 0 head 3, and 11 necessary MLP positions in layer 0. The full result is therefore better described as failed localization above the prespecified threshold than as identification of a citation mechanism. The citation-logit difference was measured at the final prompt position before generation, and the patching tables covered token positions 102–117.

To quantify the null rather than merely report an absence, we bootstrapped the mean denoising recovery of the single best-scoring component of each type over the 31 examples. No individual attention head or MLP write reached the 0.10 selection threshold on average: the best head recovered 0.025 (95% CI [0.004, 0.049]) and the best MLP position recovered 0.070 (95% CI [0.023, 0.127]). Residual-stream patches at early layers did reach the threshold, with a best recovery of 0.163 (95% CI [0.067, 0.253]) and four residual positions above 0.10 on average. The residual patches affect the selected next-token logit contrast, but no individual head or MLP write meets the prespecified positive threshold. These data do not identify a causal circuit or establish that the observed residual effect controls citations in full answers.

Validation found 7 missed-citation examples and 58 spurious-citation examples. We define correction success rate as the percentage of validation examples whose citation behavior is fixed after component scaling. For missed-citation examples, success means that the intervention adds the expected citation while preserving the answer. For spurious-citation examples, success means that the intervention removes the unwanted citation while preserving the answer. Neither group had a successful correction: 0/7 missed-citation cases and 0/58 spurious-citation cases. Because no individual head or MLP write reached the selection threshold on average, this result does not test whether scaling the residual-stream signal could control generation.

5.4. Hidden-State Probe Diagnostic

The probe datasets contained 15,328 Llama feature rows, 16,912 Qwen feature rows, and 70,272 Gemma feature rows across all extracted layers. These are not independent examples. At the selected layer, the question-level test sets contained 71 Llama rows with 1 positive label, 100 Qwen rows with 2 positives, and 226 Gemma rows with 4 positives; the corresponding train/validation positive counts were 7/2, 32/7, and 28/3. Table 5 reports performance after re-extracting every model’s features at the label-independent target-phrase anchor. Validation selected an early layer for Llama, a later layer for Qwen, and an intermediate layer for Gemma; there is no consistent depth at which the probe succeeds.

None of the probes exceeded its strongest baseline by AUROC. The point estimates are unstable because the test sets contain only one, two, and four positive examples. In 1,000 question-group bootstrap resamples, the AUROC 95% intervals were [0.147, 0.380] for Llama, [0.081, 0.961] for Qwen, and [0.073, 0.872] for Gemma; the AUPRC intervals

Model	Layer	Rel. depth	AUROC	AUPRC	Bal. acc.	Best baseline AUROC
Llama-3.1-8B	2	0.06	0.243	0.019	0.471	0.757
Qwen2.5-7B	19	0.70	0.531	0.074	0.679	0.765
Gemma-3-12B	18	0.38	0.328	0.053	0.600	0.716

Table 5

Question-level test performance of selected-layer probes using the label-independent feature anchor. Best baseline AUROC is the highest of lexical overlap, citation-logit margin, and document-query cosine similarity on the same test split.

Model	Held-out condition	Positives / rows	AUROC	AUPRC
Llama-3.1-8B	Cited other	1 / 3	1.000	1.000
	Random	0 / 34	–	–
	Relevant uncited	0 / 34	–	–
Qwen2.5-7B	Cited other	0 / 4	–	–
	Random	0 / 47	–	–
	Relevant uncited	2 / 49	0.596	0.152
Gemma-3-12B	Cited other	2 / 15	0.500	0.183
	Random	1 / 108	0.439	0.016
	Relevant uncited	1 / 103	0.833	0.056

Table 6

Leave-one-condition-out probe transfer at each model’s selected layer. A dash means that the held-out set had no positive examples, so AUROC and AUPRC are undefined.

were [0.016, 0.071], [0.011, 0.401], and [0.005, 0.319], respectively. These results do not support reliable screening of post-rationalization from the tested hidden-state features.

Table 6 reports condition transfer. Several held-out conditions contained no positive labels, making discrimination metrics undefined. The defined estimates rely on one or two positives per condition and vary substantially. Transfer therefore does not establish a condition-general faithfulness signal in these data.

6. Discussion

The primary finding is a set of operational citation shifts under answer-phrase insertion. Among recovered statements, altered random documents were cited at 0–1.1%, whereas the relevant-uncited and cited-other conditions had higher observed rates. This pattern is compatible with citation assignment responding to document-answer fit or prior citation context. It is not a causal estimate of those factors: the random condition fixes the altered document at slot [5] and sometimes replaces the original source, while the other conditions retain that source and have different eligibility rules. The event label does not reveal whether the inserted document, the original source, or stored knowledge produced the answer.

The auxiliary results set bounds on stronger claims. Most ConflictBank closed-book answers were ambiguous under the answer matcher, so the full-sample with-context match rate does not measure general override of known facts. The demonstrated-knowledge subsets are small, especially for Llama (5 examples). A context-matching answer often cites the supplied document, but the prompt instructs the model to cite it; this co-occurrence is not a provenance test. In the mechanistic diagnostic, residual-stream patches changed a selected next-token citation logit contrast, but no individual head or MLP write reached the prespecified positive threshold. Scaling selected components corrected none of 7 missed-citation or 58 spurious-citation cases. The corrected probes scored below the strongest simple baseline for every model, with only 1, 2, and 4 positives in their respective test sets. Neither diagnostic supports a deployable monitor or a validated citation mechanism.

6.1. Theoretical and Practical Implications

The distinction between textual support and causal provenance remains important for RAG evaluation. This study contributes an explicit, reproducible behavioral endpoint with separate counts for constructible interventions, recovered

statements, and citation events. The endpoint is useful as a stress test: it finds citations assigned to altered answer-bearing documents and identifies cases for closer investigation. It should not be interpreted as a measured prevalence of unfaithful citations or an isolated effect of document relevance.

For practical evaluation, claim–source support checks can be paired with controlled regeneration and source ablation. A stronger attribution study would randomize document positions, include no-insertion and placebo arms, retain the original source consistently, manually audit labels, and test whether removing a cited source changes the claim while alternative sources are controlled. Larger, independently labeled positive sets are also needed before a hidden-state screening claim can be revisited.

7. Limitations

This study uses an operational citation event under answer-string insertion. The event does not directly establish unfaithful provenance and does not cover every kind of attribution failure. The Natural Questions setup depends on BM25 retrieval over KILT Wikipedia chunks and on parser rules for cited statements. The cited-other condition has low availability for Llama, so its confidence interval is wide. The mechanistic experiment is also narrow: it uses only Llama, only PopQA capital facts, and a citation-token-versus-period logit difference as the target. The probe labels are highly imbalanced: the selected-layer test sets contain only 1, 2, and 4 positives for Llama, Qwen, and Gemma. Several condition-transfer test sets have no positive examples, and the corrected probes do not exceed the strongest baseline by AUROC. Finally, the local run used Apple MPS rather than CUDA, so exact numerical results may differ from the original CUDA-oriented setup.

The behavioral design did not include no-insertion or placebo-insertion arms, nor did it randomize document position; random forged documents were fixed at slot [5]. The reported rates therefore measure behavior under answer-string insertion and should not be interpreted as an isolated insertion effect. We also did not manually audit approximately 100 labels per model, so substring and sentence parsing errors remain possible; leave-one-condition-out probe transfer remains inconclusive because held-out conditions have zero to two positives, so condition-general prediction has not been demonstrated. The ConflictBank parametric-known subsets are small, especially for Llama (5 examples), so the corresponding override-rate interval is wide and should be read as indicative rather than precise. Finally, Gemma-3-12B is larger than the two other models, and the mechanistic sample is limited to one model and one relation type.

8. Conclusion

This paper stress-tested citation attribution in three open-weight RAG models by inserting an answer phrase into different retrieved documents. Altered documents were sometimes cited with a recovered statement, especially in the cited-other condition, but the rates are conditional on distinct eligibility and recovery funnels. The experiment measures observable citation assignment, not the causal provenance of an answer. ConflictBank answer matching, activation patching, and hidden-state probing did not supply independent validation of faithful or unfaithful source use; the latter two yielded negative generation-level and predictive results. The evidence supports a descriptive cross-model stress test and a clear evaluation boundary: citations that support a claim, citations that move under document interventions, and citations that identify the source responsible for a claim are different evidential claims. Future work needs matched document-position controls, placebo and ablation arms, manual label validation, and independent test sets before drawing stronger causal or deployment conclusions.

References

- Asai, A., Wu, Z., Wang, Y., Sil, A., Hajishirzi, H., 2024. Self-RAG: Learning to retrieve, generate, and critique through self-reflection, in: The Twelfth International Conference on Learning Representations.
- Azaria, A., Mitchell, T., 2023. The internal state of an LLM knows when it's lying, in: Findings of the Association for Computational Linguistics: EMNLP 2023, pp. 967–976. doi:10.18653/v1/2023.findings-emnlp.68.
- Belinkov, Y., 2022. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics* 48, 207–219. doi:10.1162/coli_a_00422.
- Burns, C., Ye, H., Klein, D., Steinhardt, J., 2023. Discovering latent knowledge in language models without supervision, in: The Eleventh International Conference on Learning Representations.
- van Dort, I., Heuss, M., 2026. How do LLMs cite? a mechanistic interpretation of attribution in retrieval-augmented generation, in: Proceedings of the 48th European Conference on Information Retrieval. To appear.

- Dziri, N., Rashkin, H., Linzen, T., Reitter, D., 2022. Evaluating attribution in dialogue systems: The BEGIN benchmark. *Transactions of the Association for Computational Linguistics* 10, 1066–1083. doi:10.1162/tac1_a_00506.
- Gao, T., Yen, H., Yu, J., Chen, D., 2023. Enabling large language models to generate text with citations, in: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 6465–6488. doi:10.18653/v1/2023.emnlp-main.398.
- Geva, M., Schuster, R., Berant, J., Levy, O., 2021. Transformer feed-forward layers are key-value memories, in: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Association for Computational Linguistics. pp. 5484–5495. URL: <https://aclanthology.org/2021.emnlp-main.446/>, doi:10.18653/v1/2021.emnlp-main.446.
- Google DeepMind, 2025. Gemma 3 model card. Official model documentation. URL: <https://huggingface.co/google/gemma-3-12b-it>.
- Guu, K., Lee, K., Tung, Z., Pasupat, P., Chang, M.W., 2020. REALM: Retrieval-augmented language model pre-training, in: *Proceedings of the 37th International Conference on Machine Learning*, pp. 3929–3938.
- Izard, G., Grave, E., 2021. Leveraging passage retrieval with generative models for open domain question answering, in: *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 874–880.
- Jacovi, A., Goldberg, Y., 2020. Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness?, in: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, Association for Computational Linguistics. pp. 4198–4205. URL: <https://aclanthology.org/2020.acl-main.386/>, doi:10.18653/v1/2020.acl-main.386.
- Jain, S., Wallace, B.C., 2019. Attention is not Explanation, in: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, Association for Computational Linguistics. pp. 3543–3556. URL: <https://aclanthology.org/N19-1357/>, doi:10.18653/v1/N19-1357.
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y.J., Madotto, A., Fung, P., 2023. Survey of hallucination in natural language generation. *ACM Computing Surveys* 55, 1–38. URL: <https://doi.org/10.1145/3571730>, doi:10.1145/3571730.
- Karpukhin, V., Oguz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., Yih, W.t., 2020. Dense passage retrieval for open-domain question answering, in: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Association for Computational Linguistics. pp. 6769–6781. URL: <https://aclanthology.org/2020.emnlp-main.550/>, doi:10.18653/v1/2020.emnlp-main.550.
- Kwiatkowski, T., Palomaki, J., Redfield, O., Collins, M., Parikh, A., Alberti, C., Epstein, D., Polosukhin, I., Devlin, J., Lee, K., Toutanova, K., Jones, L., Kelcey, M., Chang, M.W., Dai, A.M., Uszkoreit, J., Le, Q., Petrov, S., 2019. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics* 7, 453–466.
- Leon, M., 2025. Gpt-5 and open-weight large language models: Advances in reasoning, transparency, and control. *Information Systems* , 102620.
- Lewis, P., Perez, E., Piktus, A., Karpukhin, V., Goyal, N., Kuttler, H., Lewis, M., Yih, W.t., Rocktaschel, T., Riedel, S., Kiela, D., 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks, in: *Advances in Neural Information Processing Systems*.
- Li, K., Patel, O., Viégas, F., Pfister, H., Wattenberg, M., 2023a. Inference-time intervention: Eliciting truthful answers from a language model, in: *Advances in Neural Information Processing Systems*.
- Li, Y., Yang, N., Wang, L., Wei, F., Li, W., 2023b. Generative retrieval for conversational question answering. *Information Processing & Management* 60, 103475.
- Lin, S., Hilton, J., Evans, O., 2022. TruthfulQA: Measuring how models mimic human falsehoods, in: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics. pp. 3214–3252. doi:10.18653/v1/2022.acl-long.229.
- Liu, B., Wu, Y., Zhang, F., Liu, Y., Wang, Z., Li, C., Zhang, M., Ma, S., 2022. Query generation and buffer mechanism: Towards a better conversational agent for legal case retrieval. *Information Processing & Management* 59, 103051.
- Liu, N.F., Zhang, T., Liang, P., 2023. Evaluating verifiability in generative search engines, in: *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 7001–7025. doi:10.18653/v1/2023.findings-emnlp.467.
- Longpre, S., Perisetla, K., Chen, A., Ramesh, N., DuBois, C., Singh, S., 2021. Entity-based knowledge conflicts in question answering, in: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Association for Computational Linguistics. pp. 7052–7063. doi:10.18653/v1/2021.emnlp-main.565.
- Mallen, A., Asai, A., Zhong, V., Das, R., Khashabi, D., Hajishirzi, H., 2023. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories, in: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pp. 9802–9822.
- Maynez, J., Narayan, S., Bohnet, B., McDonald, R., 2020. On faithfulness and factuality in abstractive summarization, in: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, Association for Computational Linguistics. pp. 1906–1919. doi:10.18653/v1/2020.acl-main.173.
- Meng, K., Bau, D., Andonian, A., Belinkov, Y., 2022. Locating and editing factual associations in GPT, in: *Advances in Neural Information Processing Systems*.
- Meta AI, 2024. Llama 3.1 model card. Official model documentation. URL: <https://huggingface.co/meta-llama/Llama-3.1-8B-Instruct>.
- Nanda, N., Bloom, J., 2022. TransformerLens. Software library for mechanistic interpretability. URL: <https://github.com/TransformerLensOrg/TransformerLens>.
- Neeman, E., Aharoni, R., Honovich, O., Choshen, L., Szepietor, I., Abend, O., 2023. DisentQA: Disentangling parametric and contextual knowledge with counterfactual question answering, in: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics. pp. 10056–10070. doi:10.18653/v1/2023.acl-long.559.
- Niu, C., Wu, Y., Zhu, J., Xu, S., Shum, K., Zhong, R., Song, J., Zhang, T., 2024. RAGTruth: A hallucination corpus for developing trustworthy retrieval-augmented language models, in: *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Association for Computational Linguistics. pp. 10862–10878. URL: <https://aclanthology.org/2024.acl-long.585/>, doi:10.18653/v1/2024.acl-long.585.

- Petroni, F., Piktus, A., Fan, A., Lewis, P., Yazdani, M., De Cao, N., Thorne, J., Jernite, Y., Karpukhin, V., Maillard, J., Plachouras, V., Rocktaschel, T., Riedel, S., 2021. KILT: A benchmark for knowledge intensive language tasks, in: Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics, pp. 2523–2544.
- Petroni, F., Rocktäschel, T., Riedel, S., Lewis, P., Bakhtin, A., Wu, Y., Miller, A., 2019. Language models as knowledge bases?, in: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Association for Computational Linguistics, pp. 2463–2473. URL: <https://aclanthology.org/D19-1250/>, doi:10.18653/v1/D19-1250.
- Qi, J., Sarti, G., Fernández, R., Bisazza, A., 2024. Model internals-based answer attribution for trustworthy retrieval-augmented generation, in: Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing, pp. 6037–6053. doi:10.18653/v1/2024.emnlp-main.347.
- Qwen Team, 2024. Qwen2.5-7B-Instruct model card. Official model documentation. URL: <https://huggingface.co/Qwen/Qwen2.5-7B-Instruct>.
- Rashkin, H., Nikolaev, V., Lamm, M., Aroyo, L., Collins, M., Das, D., Petrov, S., Tomar, G.S., Turc, I., Reitter, D., 2023. Measuring attribution in natural language generation models. Computational Linguistics 49, 777–840. doi:10.1162/coli_a_00486.
- Reimers, N., Gurevych, I., 2019. Sentence-BERT: Sentence embeddings using siamese BERT-networks, in: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, pp. 3982–3992. doi:10.18653/v1/D19-1410.
- Robertson, S., Zaragoza, H., 2009. The probabilistic relevance framework: BM25 and beyond. Foundations and Trends in Information Retrieval 4, 1–174. URL: <https://doi.org/10.1561/15000000019>, doi:10.1561/15000000019.
- Su, Z., Zhang, J., Qu, X., Zhu, T., Li, Y., Sun, J., Li, J., Zhang, M., Cheng, Y., 2024. ConflictBank: A benchmark for evaluating the influence of knowledge conflicts in large language models, in: Advances in Neural Information Processing Systems.
- Turpin, M., Michael, J., Perez, E., Bowman, S., 2023. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting, in: Advances in Neural Information Processing Systems, pp. 74952–74965.
- Wallat, J., Heuss, M., de Rijke, M., Anand, A., 2025. Correctness is not faithfulness in retrieval augmented generation attributions, in: ICTIR 2025: Proceedings of the 2025 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval, Association for Computing Machinery, pp. 22–32. doi:10.1145/3731120.3744592.
- Wang, D., Tsang, S.J., Zhou, Y., 2026. Performance unfairness of large language models in cross-language fact-checking. Information Processing & Management 63, 104616. doi:10.1016/j.ipm.2026.104616.
- Wang, K., Variengien, A., Conmy, A., Shlegeris, B., Steinhardt, J., 2023. Interpretability in the wild: A circuit for indirect object identification in GPT-2 small, in: The Eleventh International Conference on Learning Representations.
- Wiegrefe, S., Pinter, Y., 2019. Attention is not not explanation, in: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Association for Computational Linguistics, pp. 11–20. URL: <https://aclanthology.org/D19-1002/>, doi:10.18653/v1/D19-1002.
- Xu, R., Qi, Z., Guo, Z., Wang, C., Wang, H., Zhang, Y., Xu, W., 2024. Knowledge conflicts for LLMs: A survey, in: Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing, Association for Computational Linguistics, pp. 8541–8565. doi:10.18653/v1/2024.emnlp-main.486.