

Highlights

- Citation correctness and citation faithfulness are evaluated separately.
- Relevant forged documents receive more citations than random documents.
- ConflictBank tests model responses when context conflicts with stored knowledge.
- Activation scaling does not control citation behavior during full generation.
- Hidden states provide a predictive signal for post-rationalization labels.