

Highlights

- A cross-model intervention benchmark separates citation correctness from citation faithfulness in open-weight retrieval-augmented generation systems.
- Cited-other interventions expose post-rationalization rates of 21.6% for Qwen, 17.8% for Gemma, and 9.1% for Llama.
- ConflictBank shows that supplied evidence shifts model answers even when it conflicts with knowledge stored in the model.
- Hidden-state probes outperform lexical, citation-logit, and semantic baselines, reaching test AUROC values between 0.831 and 0.969.
- Activation patching finds internal model components linked to citation scores, but scaling them does not fix citation errors during full generation.